

## Bubble Sort – Optimisation

### Motivation and introduction

In the last session, you have translated the C code for the *BubbleSort* algorithm to assembly code and you simulated the results with *dlxsim* and *ModelSim*. In this session, we will start to add a new instruction to the CPU to speed up the application. This new instruction should be implemented in both of *dlxsim* and in the real CPU (*ASIPMeister*). Afterwards, you will manually optimize your individual assembly code in a competition against the other groups! For every part, that starts like “a)”, “b)” ... you have to mail the answers and asked files with a CC to your group members to [asip00@ira.uka.de](mailto:asip00@ira.uka.de) and use the topic “asipXX-Session2”, with XX replaced by your group number.

### Readings for the next session: Chapters 6

### Exercises

#### 1) Adding a new instruction to *dlxsim*

1. Now we start adding a new instruction to our processor to speed up the execution. This new instruction will be “*bgeu*” (branch if greater or equal -unsigned). This instruction is going to replace the three appearances of the combination of “*sltu*” and “*beqz*”. First implement the new instruction into *dlxsim*, as explained in the Chapters 3.2.2 and 3.2.3 of the Laboratory Script (use the instruction format *BRANCH\_2\_OP* with the opcode 0xB0000000).
2. Therefore you have to copy *dlxsim* to your local home (to be able to modify it) and you have to configure the “*env\_settings*” to use your local *dlxsim* (see Figure 2-5 in the Laboratory Script).
3. Write a small assembly code to test your new instruction (also try things like:  $-4 \geq 5$ ; this should be true, as the unsigned version of  $-4$  is bigger than 5).
4. Create a new project directory inside your *ASIPMeisterProjects* directory for the new CPU and name it „*dlx\_bgeu*”. You can use a copy from the *dlx\_basis* CPU project from the last session, but do not forget to adjust the “*env\_settings*”. Also create a subdirectory for your test application in “*Applications*” directory.
5. Afterwards use this new instruction in your *BubbleSort* implementation from Exercise 4 of the last session and name the resulting file “*bs\_bgeu.s*” in a separate subdirectory in “*Applications*”.
6. Compile “*bs\_bgeu.s*” in its application subdirectory using “*make sim*” and then “*make dlxsim*”. Make sure, that the resulting array is still correct. HINT: It is ok if “*make sim*” complains that the assembler was not aware about this new instruction (and thus the binaries for *ModelSim*/*FPGA* are not created). We need to change the CPU in *ASIP Meister* to make the assembler aware of this instruction (next exercise). However, the files for *dlxsim* are created before the assembler is called.
  - a) How many cycles do you need for execution? Attach *bs\_bgeu.s* to the mail.
  - b) What is the speedup compared to “*bs\_basis.s*” (i.e. #Cycles without *bgeu* / #Cycles with *bgeu*)?

#### 2) Adding the new instruction to *ASIPMeister*

7. In your project directory start *ASIPMeister* and add the “*bgeu*” instruction to your new CPU. First you will need a new instruction format for “*bgeu*”, but the “*MicroOp Description*” will become the most difficult part (use the “*Micro-Operation Description Coding Guide*” in

“/Software/epp/ASIPmeister/share”). Have a look at “*sltu*” and “*beqz*”, as those commands are the origin for the new “*bgeu*” instruction and understand the flags of *ALU.cmpu()*. Try to change only few things compared to these instructions! Unfortunately, not everything that *should* work will produce correct VHDL code with ASIPMeister.

### Hints:

- Avoid using macros in your new *MicroOp Description*. You can test things faster if you can write them directly instead of writing them in a macro.
  - Remember, that you cannot use a hardware resource twice in the same cycle, e.g. you cannot use the ALU twice in the EXE stage. Additionally, using it in two different pipeline stage significantly complicates the whole CPU design (just think about the required wiring).
  - Remember that your number of delay-slots depends on the pipeline stage in which you write the PC. If “*bgeu*” writes the PC in another stage than the other jump instructions do, then “*bgeu*” will have a different amount of delay slots, which then would have to be considered in your assembly code.
  - You have to add a new instruction format for “*bgeu*”. Have a look at the other branch/jump instructions and see how they handle the labels.
8. Generate the hardware and software files from ASIPMeister and simulate the new instruction with ModelSim. Use the small test application that you created to test your dlxsim implementation in the previous exercises for this purpose.
  9. If everything is working fine, then simulate the *BubbleSort* assembly code that uses the “*bgeu*” instruction in ModelSim. There might be some problems if you want to use “*bgeu*” with negative values. Test it for negative values too, but if it is working fine for everything except negative values, then it is ok.
    - a) Attach the *ASIPMeister* file for the “*bgeu*” CPU to the mail.

### 3) Optimizing the assembly code

10. During the session last sessions, the *NOP* instructions have only been used to fulfil the data dependencies. This time we will try to reduce the number of *NOP*’s to reduce the total number of executed cycles. **The goal is to get the code (correctly) executed in as few cycles as possible!** This is a **competition** against the other groups and to make sure that everyone uses the same environment you have to use the framework “*bs\_Framework\_pipelined.s*”. In this framework, the main method and the remainder of the *\_bubbleSort* method have been prepared with a non-optimized usage of *NOP*’s, which is **not** allowed to be changed for your optimizations! You are also **not** allowed to remove the stack operations for e.g. saving the registers on the stack, but you are allowed to reschedule them to get rid of some *NOP*’s. Of course, you are also **not** allowed to just write the correct result into the data memory, as one could say that the initial data memory is static and so the result is always the same. Your algorithms still has to compute the correct result itself.
11. Start with creating a copy of your assembly code in a new application subdirectory with the pipelined framework and name the resulting file “*bs\_bgeu\_opt????.s*”, where the four question marks are replaced by the optimized number of cycles dlxsim needed to execute the program. Although you should work with dlxsim the most of the time, it is important to test whether your code is also running correctly in ModelSim, as this is the accurate simulation. If an assembly code is running with dlxsim but not running with ModelSim, then this is a bug in dlxsim not vice versa and thus it is a bug in your assembly code.

a) Attach your fastest-but-still-correct “*bs\_bgeu\_opt????.s*” to the mail.

**Hints:**

- For every version of your program, that is faster than the successor you should create a new “*bs\_bgeu\_opt????.s*”, as sometimes optimizations are no longer correct and then you will need a backup point from where you can start your next try.
- If you consciously use an old value of a register (with an unresolved data dependency) then you should mark this in the comment, e.g. “*r10'OLD*”. This information will become very important when you later think about further optimizations or when you search for a bug in your code.
- Carefully read and understand the examples about optimization in Chapter 3.1.1 of the Laboratory Script.
- To give you some kind of motivation and to show you what is possible: The students in the prior semesters managed to go down below **2,500** cycles to sort the array. Some of the groups were even able to break the **2,000** cycle's barrier by using some aggressive optimizations. It is not necessary to reach these numbers but you have to try to run your application as fast as possible.